

Original research paper

A comparison of methods for the analysis of binomial clustered outcomes in behavioral research



Alberto Ferrari*, Mario Comelli

University of Pavia, Department of Behavioral and Brain Sciences, Medical Statistics Section, Italy

HIGHLIGHTS

- We performed a comparison of statistical methods for the analysis of clustered binary outcomes in behavioral research with small sample sizes.
- Beta-binomial regression performed accurate and powerful hypothesis testing, outperforming even Generalized Linear Mixed Models in a range of scenarios.
- A misspecified linear model, in some circumstances, can represent a reasonable compromise between technical approachability and accuracy when dealing with proportion data.
- Poisson regression should not be applied straight away to modeling of proportion data.

ARTICLE INFO

Article history:

Received 6 July 2016

Received in revised form 5 October 2016

Accepted 7 October 2016

Available online 14 October 2016

Keywords:

Proportions

Generalized Linear Models

Generalized Linear Mixed Models

Clustered data

Beta-binomial regression

Behavioral science

ABSTRACT

Background: In behavioral research, data consisting of a per-subject proportion of “successes” and “failures” over a finite number of trials often arise. This clustered binary data are usually non-normally distributed, which can distort inference if the usual general linear model is applied and sample size is small. A number of more advanced methods is available, but they are often technically challenging and a comparative assessment of their performances in behavioral setups has not been performed.

Method: We studied the performances of some methods applicable to the analysis of proportions; namely linear regression, Poisson regression, beta-binomial regression and Generalized Linear Mixed Models (GLMMs). We report on a simulation study evaluating power and Type I error rate of these models in hypothetical scenarios met by behavioral researchers; plus, we describe results from the application of these methods on data from real experiments.

Results: Our results show that, while GLMMs are powerful instruments for the analysis of clustered binary outcomes, beta-binomial regression can outperform them in a range of scenarios. Linear regression gave results consistent with the nominal level of significance, but was overall less powerful. Poisson regression, instead, mostly led to anticonservative inference.

Comparison with existing methods: GLMMs and beta-binomial regression are generally more powerful than linear regression; yet linear regression is robust to model misspecification in some conditions, whereas Poisson regression suffers heavily from violations of the assumptions when used to model proportion data.

Conclusions: We conclude providing directions to behavioral scientists dealing with clustered binary data and small sample sizes.

© 2016 Elsevier B.V. All rights reserved.

1. Introduction

Most physiological parameters studied by biomedical researchers are continuous variables whose distribution approxi-

mates well normality; some examples of this are weight, height, blood pressure, hormone levels. For this reason, parametric methods assuming normal distribution of the response variable are the most widely used statistical instruments in biomedicine. Data showing strong departure from normality, on the other hand, are usually dealt with by transforming them to achieve better Gaussian approximation, or resorting to the use of nonparametric methods.

* Corresponding author.

E-mail addresses: alberto.ferrari04@universitadipavia.it, aferrari34@yahoo.com (A. Ferrari).

Nonparametric tests, though, suffer from decreased power and difficulty in dealing with interaction effects; these limitations suggest the use of more powerful instruments when they are available. Furthermore, in some fields of research it is not uncommon to see variables arise whose behavior, while not approximating normality, is well described by other known probability distributions.

In behavioral sciences these non-Gaussian behaviors arise naturally quite often, due to the peculiar nature of the measured responses. One such example is the outcomes of decision making tasks, in which the subject has to choose among two or more different behaviors, with one response being considered a “success” and the other(s) a “failure”; e.g. the Iowa Gambling Task (Stockard et al., 2007) or the Game Dice Task (Brand et al., 2005), used in the study of pathological gambling. In these cases, the outcome of interest is the ratio of correct choices on the number of trials.

The distributions of proportions usually do not approximate normality; they are generally asymmetrical and only admit a range of values from 0 to 1. One way to approach them is to apply the arcsine square root transformation to the data, but this has been shown to provide only minor or no improvements in power over the analysis of untransformed data, thus its use is not advised (Jaeger, 2008). Since the outcomes of behavioral experiments such as the decision making tasks we mentioned are a series of Bernoulli trials, logistic regression can be proposed as a more formal solution to their analysis. Despite this, the very common problem of overdispersion, i.e. the excess variance not accounted for by the model (usually due to overlooked sources of variation, such as inter-individual variability), can make the choice of the appropriate analytical instrument and experimental design very challenging. Simpler analytic approaches such as classical linear regression have the advantage of being very easy to apply, but at the same time must rely on their robustness to violations of some of the model's assumptions. More refined instruments are available, but they are often technically challenging even to statisticians, and their misuse can have harmful consequences. Therefore, in everyday practice, researchers must often face a tradeoff between correct model specification and technical approachability.

The aim of this paper is to review and compare some of the most relevant methods available for the analysis of proportions in the usual behavioral setup, in order to evaluate in a simulation study the robustness of hypothesis testing to model misspecification, and to provide guidelines to the reader for the choice of adequate analytic instruments and sample sizes.

2. Dealing with proportions

2.1. Linear regression and binomial model

Let us consider a behavioral experiment in which N subjects are exposed to a predetermined number n of trials, and in each trial they are required to choose between two different possible responses. Some examples of this kind of experiments are questionnaires; escape tests, in which the subject has to choose the appropriate response to avoid an aversive stimulus; or risky decision making tasks used to assess the preference of the subject for “safe” versus “risky” rewards, such as the already mentioned Iowa Gambling Task. We may want to assess whether an experimental variable, such as a genetic trait or a drug treatment, has a significant effect on the propensity of the subject towards one of the two choices.

As we noted before, the most widespread method used to deal with such results in behavioral science is linear approximation: subject becomes the statistical unit and the number of successes, or the ratio of successes to failures, is the measured outcome. As long as subject is the only grouping factor in experimental design,

observations are independent and clustering is basically removed from the picture (a more refined approach consists in regressing single outcomes on the covariates in a mixed model with subject as random effect, but this is of use mainly when within-subjects fixed effects are of interest (Stockard et al., 2007)).

This approach is not a priori unacceptable, and has a very clear advantage in practical terms, i.e. it is very easy to apply and to interpret; yet, when clustered binary data are collapsed into counts or ratios of successes, violations of the assumptions of linearity and homoscedasticity are expected and their effects on inference should be evaluated.

A more formal approach, that takes into greater account the nature of the data generating process, consists in considering each of the $n \times N$ trials as a Bernoulli process with two possible outcomes, “success” and “failure”, with probability of success π and probability of failure $1 - \pi$; in this case the number of correct responses y is a random variable with a binomial probability distribution of parameter π (Jaeger, 2008).

This approach to the data requires us to perform the analysis using Generalized Linear Models (GLMs), that allow us to model relations between the covariates and the response variable when the latter's distribution is described by a noted non-Gaussian probability function.

2.2. Poisson regression

One possible alternative to linear regression that takes more into account the data generating process is Poisson regression. Indeed, the so called “law of rare events” states that, when n is large compared to π , i.e. successes are “rare”, the binomial distribution approximates the Poisson distribution.

Poisson regression is the optimal solution to deal with count data that can be interpreted as the outcome of a binomial process with an infinite number of trials and a finite number of successes; e.g., when considering the number of occurrences of a certain event in a given amount of time (Cameron and Trivedi, 2013).

In our hypothetical experiment, the single subject is the statistical unit, and the raw number of correct choices it makes is the response to be analyzed through Poisson regression. In this case, the model will be expressed in the form:

$$\log(E(y|x)) = \beta_0 + \beta'x \quad (1)$$

Where y is the count outcome vector, β_0 is the intercept and β' the vector of fixed effect coefficients. This method is very easy to apply in most statistical software and in particular R; plus, it can be used also to model situations in which n varies from cluster to cluster in the N clusters by including an *offset* = $\log(n_i)$ into the model. Nevertheless, the efficiency of Poisson regression in a context in which n is limited is hindered by the upper bound on the number of possible correct responses, since Poisson distribution allows for all integer values in the range going from 0 to $+\infty$. Therefore, the Poisson model is also misspecified for proportion data. We can expect Poisson approximation to work well only when we have large n and comparatively low π ; in fact it has been shown that it can be a powerful alternative to linear regression even in experimental conditions where an upper bound is present (Lazic, 2015), and it has been applied to the study of complex decision making (Giang and Donmez, 2015; Paserman, 2016), gambling (James et al., 2016) and perseverative behavior (Lazic, 2015).

Another issue the experimenter might meet when applying Poisson regression is the inflation of Type I error rate in presence of overdispersion; indeed, inference in Poisson regression is heavily dependent on the assumption of equality of mean and variance. However, this problem can be fixed by applying robust sandwich

standard error estimators to the model (the so called Huber-White estimator) (Silva and Tenreiro, 2006; White, 1980).

2.3. Beta-binomial regression

Beta-binomial regression has been proposed as an alternative to linear regression to model clustered binary data, such as in the case of the proportion of successes and failures over a definite number of trials. This model assumes that the response variable follows a beta-binomial distribution, in which overdispersed binomial data are handled by letting the π parameter of the binomial vary randomly, following a beta distribution. The resulting distribution has probability mass function:

$$\binom{n}{y} \frac{B(y + \alpha, n - y + \beta)}{B(\alpha, \beta)} \quad (2)$$

Where y is the binary outcome vector, α and β are shape parameters and B is the beta function:

$$B(x, y) = \int_0^1 t^{x-1} (1-t)^{y-1} dt \quad (3)$$

The advantages of this approach compared to ordinary logistic regression are self-evident: overdispersion is naturally accounted for by the variability in the underlying beta distribution. Beta-binomial regression is also advantaged compared to the simpler beta regression of the observed frequencies, in that it can accommodate observed frequencies of 0 or 1, which may be easily arise in scenarios in which n is small.

The use of beta-binomial regression in behavioral and biomedical research in general has been for some time outside of most experimenter's reach because of lack of appropriate software implementation, but is now easily accessible using various statistical software, e.g. the R packages *aod* (Lesnoff and Lancelot, 2012) or *aods3* (Lesnoff and Lancelot, 2013), and its suitability to the analysis of proportions has been assessed previously in various contexts (Crowder, 1978; Hilbe, 2013; Muniz-Terrera et al., 2016).

2.4. Generalized Linear Mixed Models

The fifth available method we take into account are Generalized Linear Mixed Models (GLMM₅); in particular, we will consider mixed effects logistic regression with random intercepts.

The most correct way to deal with binary outcomes when the assumption of independence of the observations is met is the ordinary logistic model. This consists in modeling the response variable as the logarithm of the ratio between the probability of success π and the probability of failure $1 - \pi$; the *log-odds* or *logit*, as follows:

$$\text{logit}(E(y|x)) = \log \frac{\pi}{1-\pi} = \beta_0 + \beta'x \quad (4)$$

However, this cannot be applied straight away in presence of clusters of more or less highly correlated observations, such as in the case of repeated measures on the same subject. One way to fix this issue is using GLMMs. This allows us to model the inter-cluster variability by fitting a mixed effect model with “cluster” as random effect and a series of n_i binary outcomes for each i^{th} cluster in $i = 1, \dots, N$. In our hypothetical experiment, we will have N subjects with n_i trials for each i^{th} subjects. Since we expect the observations made on the same subject to be correlated, we allow the model to fit a different intercept for each i^{th} subject.

In this case, the model will have the form:

$$\text{logit}(E(y|x)) = \log \frac{\pi}{1-\pi} = \beta_0 + \beta'x + b'z \quad (5)$$

Where b' is the vector of coefficients for the random effects and z is the vector of the cluster-specific random effects. This is perhaps the most formally rigorous analytic method to model the kind of data we described (Jaeger, 2008), but it is also somewhat complex to approach. In particular, if:

$$b \sim_{i.i.d.} N(0, \sigma_b^2)$$

Then the likelihood for a model of this form is:

$$L = \int \prod_{i,j} f_{y_{i,j}|z_i}(y_{i,j}|x_{i,j}, b_i) f_{b_i}(b_i) db = \prod_i \int_{-\infty}^{+\infty} \frac{\exp\{\beta_0 + \beta'x_{i,j} + b_i z_i\}}{\sum_j \exp\{\beta_0 + \beta'x_{i,j} + b_i z_i\}} \frac{\exp\{2\sigma_b^2\}}{\sqrt{2\pi\sigma_b^2}} db_i \quad (6)$$

And its maximization requires dealing with intractable multi-dimensional integration, which can only be approached through numerical approximation. There are various methods and software packages for GLMMs fit able to deal with this issue, and the choice of a method over another can affect inference significantly see Bolker et al. (2009) for review. Here we will concentrate only on two of the most widely used estimation techniques: Penalized Quasi-Likelihood (PQL) and Laplace method.

PQL is probably the most commonly utilized instrument for GLMM fitting. The theoretical approach behind PQL consists in replacing the Likelihood function with a Quasi-Likelihood function which shares some properties with the true Likelihood, which is then integrated using Laplace approximation and maximized (Breslow and Clayton, 1993).

An alternative is to approximate directly the solution to the integrated Likelihood Function with Laplace method (Raudenbush et al., 2000). This is more computationally intensive than PQL and is more prone to give issues of numerical stability, but it can theoretically give more accurate results (Breslow, 2004). Both methods are implemented in advanced statistical software and in R.

In the following sections, beside using simulated data to assess the efficiency and robustness of the different analytical instruments we have listed, we will also compare the performances of these two distinct methods for GLMM fit.

3. Materials and methods

3.1. Simulation study

In order to assess power and Type I Error rate of the different methods we listed, on inference on fixed effects, we generated hypothetical data from a random intercept logistic model (5) under the following conditions:

- Three sample sizes: 16, 24, 32;
- Random intercepts generated from a centered Gaussian distribution with $\sigma = 0.5$ or 1.75 , representing “weak” and “strong” overdispersion;
- Four cluster sizes: 30, 15, 8, and varying from cluster to cluster with $\min = 1$ and $\max = 30$;
- Two fixed factors, “factor 1” and “factor 2”, with the coefficient for factor 1, β_1 , ranging from 0 to 3 (or saturation of statistical power), and the coefficient for factor 2, β_2 , fixed to 0, and the covariates assuming value 0 or 1.

Cluster number and size have been chosen as they are reasonable conditions in usual behavioral experiment designs. The two σ for the distribution of random intercepts were chosen since they give rise, respectively, to a bell-shaped approximate normal distribution and to an almost-uniform distribution over the $[0, 1]$ interval for the probability of correct response π . Datasets with varying

cluster size were generated by assigning to each cluster of size 30 a different probability of non-response; the probability of non response was obtained as the logit transformation of a Gaussian random variable with mean 0 and standard deviation 0.5 which translates into an overall mean probability of non-response of 0.5. Factor 2 was introduced in the models in order to simulate the effect of an ineffective treatment on hypothesis testing. The values of β_1 were chosen so that a very vast range of effect sizes is taken into consideration; with covariate values being 0 or 1 and β_1 ranging from 0 to 3, the values of π ranged from 0.5 to ≈ 0.95 , an interval which includes very extreme conditions in which model misspecification is expected to have the heaviest consequences.

We simulated 1000 datasets for each combination of parameters.

A mixed effect logistic model with cluster as random effect and a beta-binomial regression model were fitted to the raw generated samples. Records from each cluster were collapsed into counts to be analyzed with Poisson regression and linear regression. In scenarios with non-constant cluster size records were collapsed into proportions of *correct/incorrect* responses for linear regression. In Poisson regression with non-constant cluster size an offset equal to $\log(n)$ was included into the model.

Random effects bound to clusters are handled differently in each model. In GLMMs a random intercept term is added accounting for cluster effect; in beta-binomial regression it is modeled by a multiplicative overdispersion parameter; in linear and Poisson

regression the random effect is bypassed by making cluster the statistical unit.

We tested the hypothesis of the fixed effect coefficients β_1 and β_2 being different from 0 by using Wald's test to compute p-values, and fixing the significance level at $\alpha = 0.05$; in mixed logistic regression degrees of freedom were calculated using Satterthwaite approximation; in Poisson regression, both ordinary and robust standard error estimators were used for calculation of p-values. P-values significant at $\alpha = 0.05$ were counted to generate power curves; we also calculated Type I Error rates by counting significant p-values when both coefficients were = 0. When the models showed high rates of non-convergence or computational errors (>10%) the scenario was discarded. The results of the analysis are reported in Figs. 1 and 2 and Table 1.

The entire analysis was performed on R (RCoreDevelopmentTeam, 2014), in particular using packages MASS (Venables and Ripley, 2002), lme4 (Bates et al., 2014) and aod (Lesnoff and Lancelot, 2012).

3.2. Example 1: escape test

The first example we will deal with is the avoidance test for Escape Deficit (ED), an instrument used for the evaluation of learned helplessness in animal models, assessing the ability of the subject to develop an avoiding response to repeated aversive stimuli (Krishnan and Nestler, 2008).

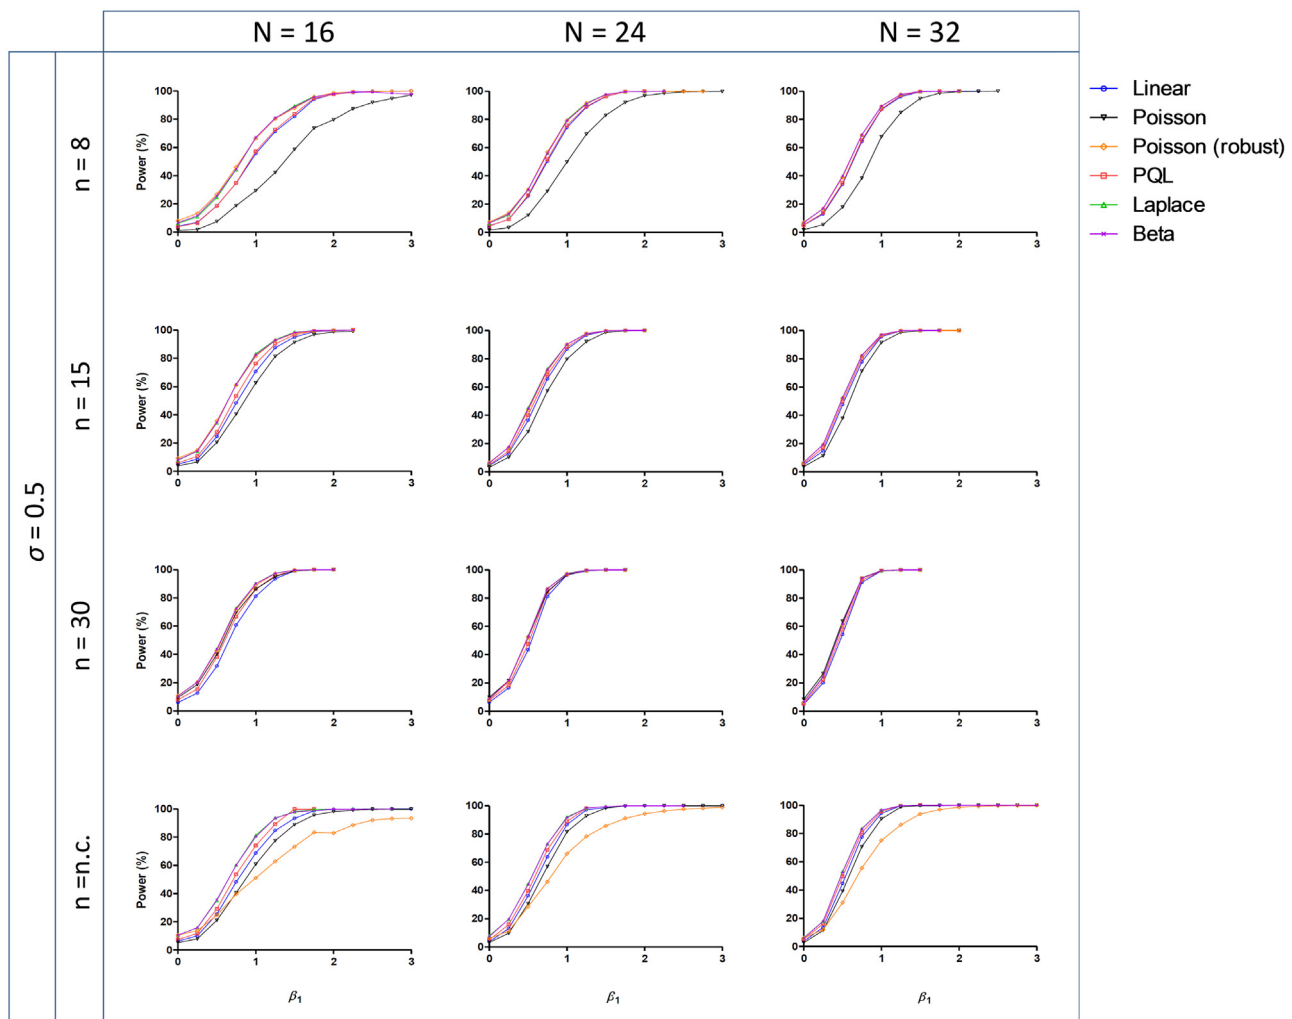


Fig. 1. Power curves over 1000 simulations for the twelve different combinations of sample size and cluster size with “weak” overdispersion.

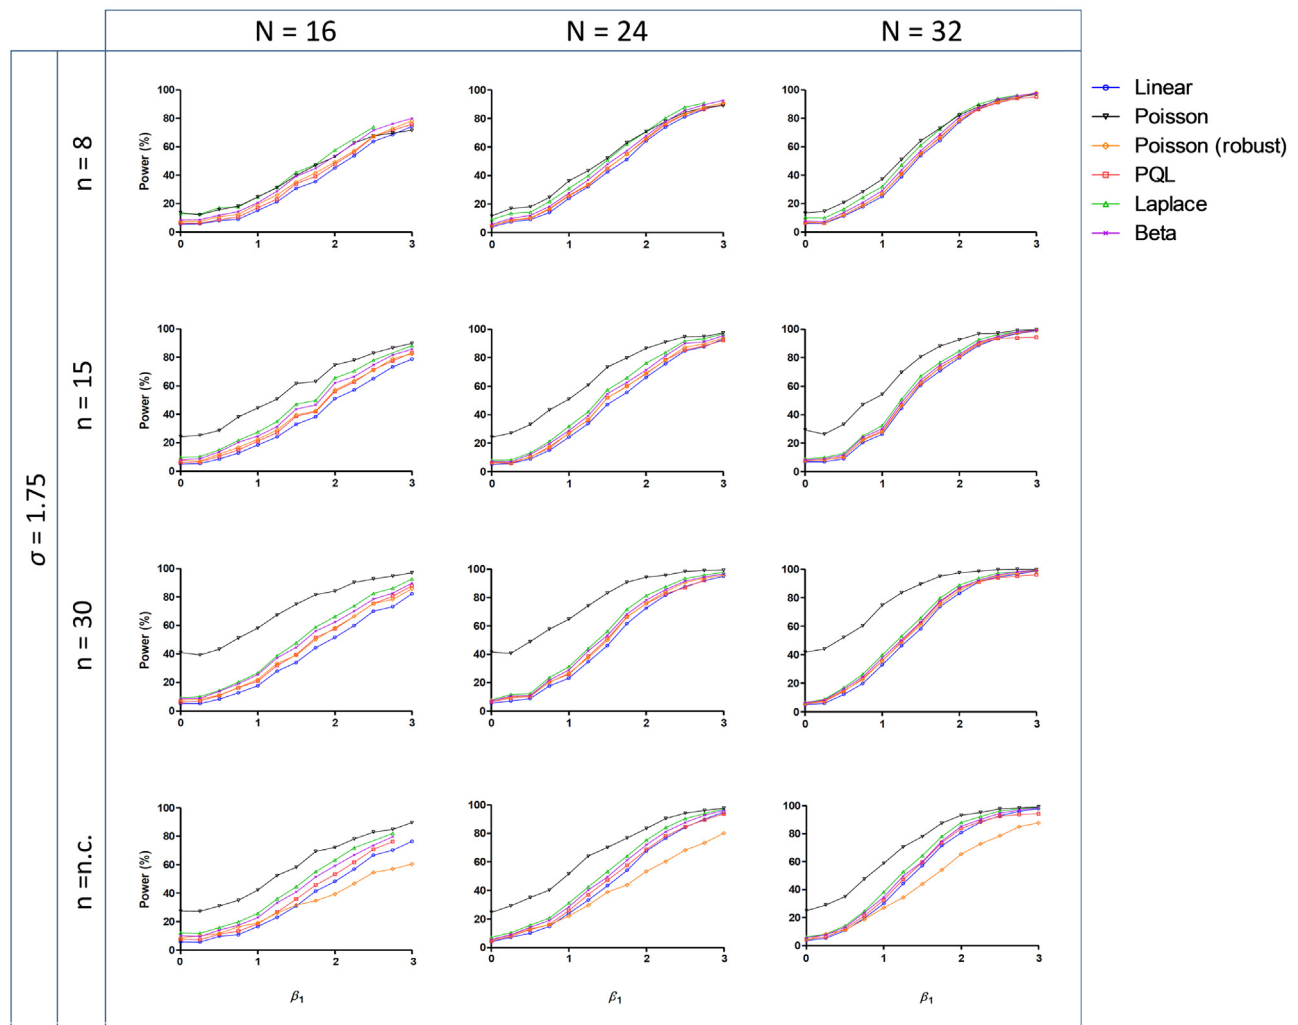


Fig. 2. Power curves over 1000 simulations for the twelve different combinations of sample size and cluster size with “strong” overdispersion. Missing points are due to exceedingly high percentages of computational errors in model fit.

In this test the animal is placed in an apparatus consisting of a cage with dark walls, divided into two equal chambers by a dark partition with a sliding door; one half of the cage is connected to a generator able to deliver weak electric shocks to the animal (the electric chamber), while the other is not (neutral chamber, Fig. 3A).

After an habituation period in the apparatus, the animal is placed in the electrified chamber and receives a series of few-seconds long electric shocks, usually at a 30 s intervals from each other, in coincidence with the opening of the door connecting the electrified chamber to the neutral one. In normal conditions, the animals easily learn to avoid the aversive stimulus most of the times by moving to the neutral chamber when the shock is delivered. After each trial, animals succeeding in escaping are gently placed again in the electrified chamber and the procedure is repeated a predetermined number of times.

However, it has been shown that the ability to learn the avoidant behavior is disrupted if the animal is previously subjected to a stress protocol involving administration of inescapable electric shocks inducing learned helplessness. This effect can be reverted to various extents by treatment with appropriate psychoactive drugs, such as antidepressants or some mood stabilizers (Gambarana et al., 2001).

In this kind of experiment, we have a dichotomous outcome (escape/failure) which is repeated n times for each of the N animals, and we want to assess whether animals administered with different treatments show a significantly different performance in this

task. We reanalyzed a dataset from a study by Scheggi et al. (2015), coming from an ED test in which rats ($N = 17$) were assigned to three experimental groups and tested in the ED paradigm. 6 animals were subjected to a chronic stress protocol prior to the experiment, 6 animals were subjected to stress and treated with a classic antidepressant drug (Imipramine), and 5 animals were treated with saline solution and acted as controls. The stress protocol utilized involves administration of a series of inescapable electric shocks, aimed at inducing learned helplessness (Gambarana et al., 2001). The procedure lasted 42 days: on the first day rats were restrained in a flexible wire-net for 50 min, during that time they received a series of about 80 unavoidable electric shocks; on day 2 they were tested for escape deficit for the first time. From day 3 to day 42, a chronic stress protocol was applied that involved administration of minor stressors (10–20 min restraint with or without electric shocks) every other day; this procedure ensured that the escape deficit was maintained for a prolonged time window, during which different behavioral procedures and drug treatments were performed. Starting from day 21 Imipramine treatment was administered until the end of the experiment. On day 42, a second test for escape deficit was performed, aimed at assessing learned helplessness and the efficacy of Imipramine in reverting it; the data used in this paper come from this second test session. During the ED tests each animal received 30 electric shocks, and the responses, escapes or failures, were registered.

Table 1
Type I error rates in the simulations; 'n.c.' = non-constant.

Type I Error rate (%)		β_1				β_2			
Coefficient									
Cluster size		n = 8	n = 15	n = 30	n.c.	n = 8	n = 15	n = 30	n.c.
N = 16									
Linear	$\sigma = 0.5$	4.2	5.1	5.9	6.3	4.2	4.8	7.2	5
PQL		3.8	6	7.8	7.5	3.8	5.5	8.0	5.6
Laplace		6.2	7.8	10.6	10.6	6.2	8.6	9.9	9.1
Poisson		1.3	3.9	9.2	5.3	1.1	3.9	9.4	3.7
Poisson (robust)		8	9.2	10.6	10.5	7.3	9.2	10.4	8.9
Beta		6.6	8.1	10.4	10.5	6.4	9.3	10	8.9
$\sigma = 1.75$									
Linear		5.5	5.1	5.1	5.8	4.9	4.8	4	6.4
PQL		6.3	6.1	6.2	7.7	5.6	5.7	7.2	7.1
Laplace		13	9.7	9	12.1	12.6	10	9.5	11.6
Poisson		13.8	24.3	41.1	27.4	13	28.4	41.9	25.8
Poisson (robust)		7.6	7.7	7.6	8.6	7	7.3	7.8	8.5
Beta		8.7	8	8.5	10.2	8.5	8	8.8	10.2
N = 24									
Linear	$\sigma = 0.5$	4.6	4.3	6.1	3.8	5.6	6.2	4.6	5.9
PQL		4.3	6	7.6	5.1	5.4	5.5	6.2	7.4
Laplace		6.5	6.2	8.5	7.9	6.7	8.9	8	8.7
Poisson		1.5	3.2	9.7	3.2	1.7	3.9	7.5	5.3
Poisson (robust)		7.2	6.5	8.3	6.2	7.1	9	7.7	8.5
Beta		6.6	6.2	8.5	7.4	6.7	8.8	8	8.7
$\sigma = 1.75$									
Linear		3.9	5.1	5.6	4	4.6	4.2	6.8	4.7
PQL		4.7	6.2	6.5	4.5	5.1	5.3	7.9	5.8
Laplace		9	8	7.9	7.2	9.8	7.8	10.6	8.1
Poisson		11.5	24.3	41.5	24.7	12.1	25.5	43	26.3
Poisson (robust)		4.9	6.7	7.3	5.4	5.9	5.5	9	7.1
Beta		5.8	7.3	7	5.8	6.5	6.1	9.2	7.2
N = 32									
Linear	$\sigma = 0.5$	5.1	4.9	4.7	3.7	4.1	5.1	5.1	3.8
PQL		5.2	5.6	6.5	5.1	4.2	5.8	7.8	5
Laplace		7.1	6.6	5.3	6.2	5.5	6.9	6.4	7
Poisson		1.8	3.5	8.4	2.8	1.8	4.2	8.5	3.8
Poisson (robust)		7	6.5	6.3	6	5.9	6.4	7.9	6.2
Beta		6.8	6.7	6.4	5.9	5.7	6.7	7.9	6.3
$\sigma = 1.75$									
Linear		6.3	6.9	4	3.5	6.4	5.3	3.9	6
PQL		6.6	7.9	5.4	4.5	6.6	5.9	4.6	6.7
Laplace		10.2	9	6.4	6.2	10.2	7.2	5.7	8.3
Poisson		13.4	29.3	41.8	24.9	16.2	24.4	40.2	25.8
Poisson (robust)		7.2	7.6	5.6	4.2	7.3	6.2	4.8	6.9
Beta		7.9	8.2	6.3	5.7	8	6.7	5.2	7.7

In each model experimental group was treated as a categorical covariate with three levels: Control, Stress, Stress + Imipramine. We tested the hypothesis that animals in the Stress and Stress + Imipramine groups have an ability to enact the avoiding response significantly different from the one of control animals. To this aim we used linear regression of counts on group, Poisson regression of counts on group (with and without robust standard error estimation), beta-binomial regression and a GLMM with subject as random effect. The results are reported in Fig. 3B and Table 2.

3.3. Example 2: preference test

The data for our second example come from a study from Marchese et al. (2013) about the effect of stress and long-term treatment with lithium on the acquisition of operant behavior in rats. In a similar fashion as in example 1, here experimental subjects were assigned to three groups (controls, stressed, and stressed treated with lithium, for a total $N = 6 + 6 + 8 = 20$), and tested for their competence to acquire a vanilla sugar (VS)-reinforced instrumen-

Table 2
P-values for the treatment effects of Stress and Stress + Imipramine on the odds of correct response in the ED experiment, as calculated from the models fitted. Significance levels compared with control group: 0.1, • (suspect); 0.05, * (significant); 0.01, ** (highly significant); 0.001, *** (decisive).

p-values		Stress vs Control		Stress + Imipramine vs Control	
Linear		0.0057	**	0.1285	
GLMM	PQL	0.0058	**	0.2237	
	Laplace	0.0013	**	0.2005	
Poisson	MLE	$8.75 \cdot 10^{-9}$	***	0.0086	**
	Robust	0.0037	**	0.0902	•
Beta-binomial		0.0009	***	0.0825	•

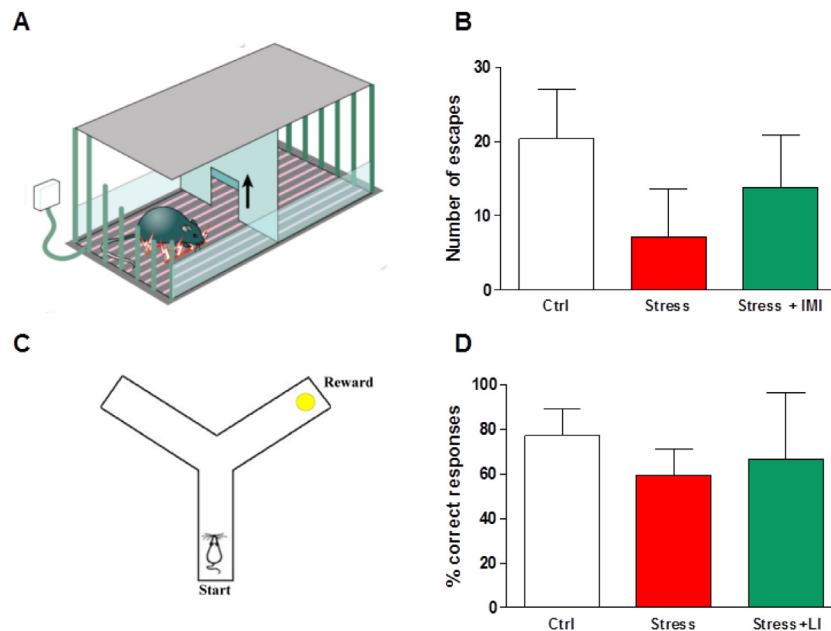


Fig. 3. A: image illustrating the setup of the avoidance learning test (adapted from Krishnan and Nestler, 2008 with permission); B: number of escapes per experimental group in the ED experiment from Scheggi et al. (2015). A difference between the stress alone and the control groups is apparent. Results are presented as mean \pm SD; C: image representing the setup for the Y-maze preference test; D: percentages of correct responses in the VAB experiment from Marchese et al. (2013). No remarkable difference can be seen among the three groups. Results are presented as mean \pm SD.

tal behavior [VS-sustained appetitive behavior (VAB)] in a Y-maze preference test paradigm. In each trial, one animal was placed in the starting arm of an Y-maze ($15 \times 40 \times 20$ cm for each arm, Fig. 3C) with a reward (a pellet of vanilla sugar) in one of the divergent arms. Each subject underwent ten daily trials, and in each trial one of three possible responses was recorded: *correct response*, if the animal moved into the arm containing the reward; *incorrect response*, if the animal went into the empty arm; *incomplete trial*, if the animal stayed in the starting arm until the end of the trial.

In the original work, the experiment was repeated daily for ten days in order to study the ability of the animals to learn the instrumental behavior, and the variation in time of the number of correct responses, incorrect responses and incomplete trials was analyzed with repeated measures ANOVA. Yet, this kind of design includes two potential levels of clustering of the binary outcomes (subjects and days), which goes beyond the scope of our study; thus, we only analyzed the outcomes from day 2. This was chosen since in the study by Marchese et al. differences between the groups were not yet detectable at this point in time. Here we did not look into the number of incomplete trials, but only into the proportion of *correct/incorrect* responses.

In each model experimental group was treated as a categorical covariate with three levels: Control, Stress, Stress + Lithium. We tested the hypothesis that the stress and stress + lithium groups differ from the control group in probability of correct response. The following models were fit to the data: linear regression of proportion of correct answers on group, Poisson regression of the number of correct answer on group (an offset equal to $\log(n)$ was added to the model, both maximum likelihood and robust standard error estimators were applied), beta-binomial regression, a GLMM with subject as random effect.

Results of the analysis are presented in Fig. 3D and Table 3.

4. Results and discussion

Even though the non-normality of the distribution of proportion data may not be an issue when dealing with a sizeable number of observations, such as in the case of, e.g., phase III clinical tri-

als, it is likely to become a serious confound in conditions where the number of observation is more limited. This is typically the case for behavioral studies in neuropsychology or cognitive science, that often rely on small samples only adequate to the detection of large effect sizes (Marszalek et al., 2011). Small sample sizes become a compelling issue, also, in preclinical drug research on animals: in this case, ethical and economic concerns alike call for the minimization of the number of subjects used in the experiments; furthermore, in experiments where the repeated trials are a source of stress for the subject (e.g., the escape tests), it is in the experimenter's interest to minimize confounding effects due to excessive stress. In these conditions, the choice of appropriate analytic instruments, sample size and number of trials are pivotal for the study to answer its experimental questions; on the other hand, the usual methods for sample size calculation are not entirely adequate to deal with non-normally distributed clustered data.

One advisable alternative is to perform preliminary power analysis through simulation (Johnson et al., 2015; Kain et al., 2015), but this is time consuming, technically challenging, and it must be tried on a wide predetermined range of scenarios in order to determine a power curve. For this reason, in this work we wish to provide

Table 3

P-values for the treatment effects of Stress and Stress + Lithium on the odds of correct response in the VAB experiment, as calculated from the models fitted. Significance levels compared with control group: 0.1, • (suspect); 0.05, * (significant); 0.01, ** (highly significant); 0.001, *** (decisive).

p-values		Stress vs Control	Stress + Lithium vs Control
Linear		0.194	0.380
GLMM	PQL	0.200	0.825
	Laplace	0.176	0.819
Poisson	MLE	0.447	0.911
	Robust	0.039	0.759
Beta-binomial		0.204	0.745

some practical guidelines to the choice of the appropriate statistical instruments and experimental design, in some scenarios that can reasonably be met by behavioral researchers dealing with heavy limitations on sample and cluster size.

Among the statistical methods taken into consideration in our study, linear approximation performed surprisingly well on proportion data, showing a Type I error rate consistent with the nominal 5% level of significance. Nevertheless, it was also more conservative compared with GLMMs and beta-binomial regression. As expected, this gap in power became less and less significant as the number of observations increased, and at $N=24$ and $n=30$ it was quite marginal. The unremarkable gain in power granted by the use of GLMMs or beta-binomial regression when both N and n are big enough, considering the relatively challenging nature of these methods, may advise against their use when the above conditions are met.

Poisson regression performed quite poorly in our study, both applying ordinary and robust standard errors estimates. This was particularly evident when we applied ordinary Poisson regression to datasets with strong overdispersion ($\sigma = 1.75$), where we found Type I error rates as high as 43%. The issue was easily fixed when sandwich standard error estimators were applied, but this came at the cost of a remarkable loss in power compared to other methods. We therefore conclude that Poisson regression, even with robust standard error estimators, should not be seen as a first choice when dealing with proportion data, and that better alternatives are available even when cluster size is constant but moderate.

For what concerns logistic GLMMs, a noticeably different outcome between fit performed through PQL and Laplace approximation was found. PQL was generally more powerful compared to linear approximation, but not to beta-binomial regression. Nevertheless, the gain in power compared to linear regression was quite thin, being as small as 0.2% in some conditions, namely large sample size, small cluster size and weak overdispersion. This was balanced by a very acceptable Type I error rate, spanning from a minimum of 3.8% to a maximum 8.0%.

When parameter estimation was performed through Laplace approximation, the logistic mixed model was generally the most powerful method, but it also had the highest Type I error rates in our experiment (with the exception of simple Poisson regression, which also gave unacceptable Type I error rates). Also, in some scenarios, namely small or non-constant n , small N , strong overdispersion and high β_1 , this method was prone to give serious issues of non-convergence. PQL, while being overall less powerful than Laplace, was somewhat less prone to this kind of problems. This is consistent with previous findings from simulation studies of multilevel logistic models comparing different methods for GLMM fit (Kim et al., 2013; Moineddin et al., 2007).

In this respect, Moineddin et al. (2007) and Kim et al. (2013) studied the performances of different methods for logistic GLMM fit in parameter and standard error estimation. We, on our part, did not look into properties of parameter estimation *per se*, focusing on hypothesis testing instead. This choice was dictated by the demands of preclinical research with limited sample sizes, that rarely allow for the detection of very small effect sizes, therefore making the issue of detectability a priority.

Beta-binomial regression, on its part, was shown to have some very interesting properties. While having a slightly higher Type I error rate compared to linear regression and PQL, it was also consistently more powerful; this was particularly evident in “worst case scenarios”, with small or non-constant n , small N and strong overdispersion. Also, it had type I error rates lower than Laplace fit, while at the same time being more robust to computational errors.

To date, beta-binomial regression has found use in a series of biomedical application, e.g. in differential methylation analysis or

in the study of cognitive function in aging research (Robinson et al., 2014; van den Hout and Muniz-Terrera, 2016). Its flexibility and robustness advise in favor of a broader use of the method in neuropsychology.

For what concerns sample and cluster size, we observed that in presence of weak overdispersion increasing sample size had similar effects on power to increasing cluster size. The most relevant improvement in this sense was observed in the shift between $N=16$ and $N=24$, and between $n=8$ and $n=15$. On the other hand, as expected, in presence of strong overdispersion changes in cluster size had little effect on power, whereas increasing sample size had more relevant consequences; e.g., using beta-binomial regression and a $\beta_1 = 1.5$, there was only an increase of roughly 4% points in power by increasing n from 8 to 15, whereas the improvement was of about 8 points when increasing N from 16 to 24.

We also compared the performances of the models on two real datasets from actual behavioral experiments on rodents, namely escape test and Y-maze preference test. In the first case, we had a small sample size of 17 subjects and a large cluster size of 30 trials, in the second case we had a moderate sample size of 20 subjects and a non-constant cluster size with a maximum of 10 completed trials. Results from example 1 are mostly consistent with simulation outcomes; in fact, the mixed logistic model fit through Laplace method and beta-binomial regression gave the lowest p-values. Poisson regression gave even lower p-values, but this result was likely a false positive due to unaccounted overdispersion. Indeed, we ran a test for overdispersion on the Poisson model from example 1 using the *dispersiontest* function from R package *AER* (Kleiber and Zeileis, 2008), discovering that there was indeed strong evidence for overdispersion (p-value = 0.0037).

Some of the results from these analyses seemed, at first sight, inconsistent with the outcomes of simulation; in particular, in the VAB experiment Poisson regression with sandwich estimators was the only method able to detect a difference between groups. This may seem puzzling, since in simulated scenarios with non-constant cluster size robust Poisson regression was not the most liberal analytic method. Yet, it has to be remembered that the Poisson model is misspecified for proportion data; thus, the overdispersion in the logit model generating the data does not necessarily show up as such when a Poisson model is applied. In fact, the test for underdispersion was highly significant in the Poisson model from example 2 (p-value = 6.732×10^{-16}). This is not unexpected since Poisson regression is expected to show its limits when n is not very big compared to π . In particular, the upper bound to the number of responses is likely to put a strong constraint on variance, resulting in underdispersion and conservative inference. This underdispersion is accounted for when robust standard error estimators are applied.

In fact, plotting the ratio between residual deviance and residual degrees of freedom from the Poisson models applied to our simulated datasets, we find that, when the logit model generating the data is only weakly overdispersed and cluster size is small, application of a Poisson model frequently leads to underdispersion, whereas when the logit model shows strong overdispersion so does the Poisson model (Fig. 4).

Indeed, in our simulated scenarios with non-constant cluster size, maximum cluster size was 30, whereas in the VAB experiment it was only 10. Also, if we assume a mixed effect logistic model as the data generating process, the estimate for β_1 obtained using Laplace method to approximate maximum likelihood is -0.7239 . In our simulations, under weak overdispersion, small N and n , and $\beta_1 < 1$, robust Poisson regression was the most liberal method, slightly exceeding in this respect even beta-binomial regression and GLMMs.

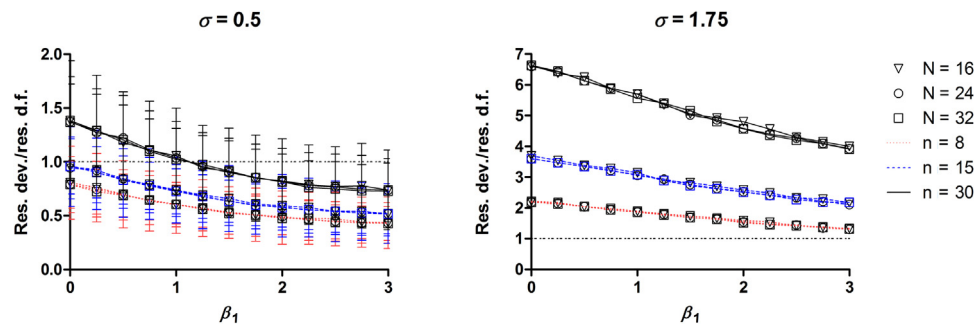


Fig. 4. Residual deviance to residual degrees of freedom ratio in Poisson models as a function of fixed effect coefficient, sample and cluster size, and degree of overdispersion of the logit model; the horizontal interrupted line represents ideal equidispersion. Data are presented as mean \pm SD over 1000 simulations.

5. Conclusions

Our study has the obvious constraints of a simulation study, plus a few other limitations.

Firstly, we only studied the performance of random intercept models, thus, our results are not directly applicable to scenarios in which single subjects respond differently to a certain treatment or condition. Another limit of our work is that we only took into consideration one source of overdispersion, individual clustering, thus assuming that each individual subject is independent of the others. In behavioral sciences, though, especially when working on animal behavior, other sources of extra variability often appear, e.g. litter and cage effect. We did not look into the effects of multiple nested and/or crossed levels of clustering, nevertheless these are likely to be relevant in the everyday practice of the experimenter. Furthermore, we did not simulate scenarios including interaction effects. The parameters of the simulations, on the other hand, cover a vast range of effect sizes and inter-individual variability; indeed, some situations were so “extreme” that in some scenarios attempts to fit appropriate binomial or beta-binomial models failed in most simulations.

The repertoire of methods for the analysis of proportion data that we took into consideration in the present study must not be considered in any way as complete or definitive. Different methods and software packages for GLMM fit exist that we have not put on trial, such as Monte Carlo Expectation Maximization (Levine and Casella, 2001) or Gauss-Hermite Quadrature (AGQ) (Clarkson and Zhan, 2002). AGQ in particular is relevant since it is purportedly more accurate than PQL and Laplace, but it was neglected since it is more computationally intensive than Laplace approximation while using basically the same principle. Indeed, Laplace method uses a Taylor expansion of the likelihood function around one point, while AGQ uses multiple points; therefore, AGQ is in a sense a slower but more accurate version of Laplace (Clarkson and Zhan, 2002). In particular, Kim et al. already showed that it performs better than Laplace with small sample sizes, but the difference tends to vanish as sample size grows (Kim et al., 2013).

Another limitations of our study is our choice to apply only fully parametric models, thus neglecting semi-parametric methods for modeling clustered data, such as Generalized Estimating Equations (GEE).

Also, most importantly, we chose to keep a frequentist framework based on classical hypothesis testing, and did not take into account Bayesian methods. Yet, these represent powerful alternative instruments for the analysis of clustered binary responses.

Despite these limitations, our study provides some directions to behaviorists and neuropsychologists working with proportion data and relatively small sample sizes:

- If we expect little to moderate inter-individual variability, the experimenter can change both sample size and cluster size to attain greater power. In conditions where the single trials are stressful or tiresome to the subject, or there is some other reason to minimize the number of trials, a reasonable compromise between sample size and cluster size for a two-factors design may be using about 15 trials on 24 subjects. In case of a significant probability of non-responses, this has to be adjusted in consideration of the expected number of *completed* trials.
- If strong inter-individual variability is expected, it is advisable to have a bigger sample size; 24 subjects are a reasonable choice for a two-factors design. The number of trials, instead, becomes less important, although it is advisable to have at least about 15 completed trials per subject.
- GLMM fit through Laplace approximation was the most powerful method in most scenarios, while being also prone to Type I errors. Our study confirms that, if the clustering structure is well known, the logistic mixed model can be seen as a good first choice for analyzing proportion data. This is expected to hold true also when AGQ is used for maximum likelihood estimates.
- GLMM fit through PQL was slightly more powerful than linear regression, and had a Type 1 error rate mostly consistent with the nominal significance level. For these reasons, though it is apparently less powerful than fit through Laplace method, it is a good first choice for analysis of proportion data.
- In our simulations, beta-binomial regression showed the most interesting properties. In terms of power, it had a performance very similar to GLMM fit through Laplace method, but at the same time had distinct advantages in terms of robustness to computational errors and of Type I error rate. Therefore, it is a very powerful instrument to deal with clustered binomial data and, depending on experimental conditions, it can even be preferable to GLMMs.
- Overall, linear regression can be an acceptable method to perform hypothesis testing with proportion data, at least for some clustering structures and experimental designs, being quite robust to model misspecification. It is in general more conservative than mixed binomial or beta-binomial models, thus it may not be the best choice for the detection of small effect sizes. Also, caution should be exercised in the interpretation of the estimated parameters, since using linear models to deal with proportion data can give rise to nonsensical predictions, such as probabilities outside the [0,1] interval.
- Poisson regression should not be applied straight away to proportion data because of the strong inflation of Type I error rate it can produce when strong overdispersion is present. Robust Poisson regression, on the other hand, has a reasonable Type I error rate but is also not very powerful compared to beta-binomial regression or GLMM fit through Laplace method. Yet, it may be

advantaged in the detection of small effect sizes when weak overdispersion is expected. This approach requires great caution, though, since with a small number of trials the Poisson model is heavily misspecified, and will likely give rise to biased estimates and nonsensical predictions. In any case, testing for overdispersion and/or underdispersion is essential when applying Poisson regression to proportion data.

Acknowledgements

We wish to thank Maria Graziella De Montis, Carla Gambarana and Simona Scheggi of the University of Siena for providing data used in the examples.

References

- Bates, D., Mächler, M., Bolker, B., Walker, S., 2014. Fitting linear mixed-effects models using lme4. arXiv preprint arXiv 1406.5823.
- Bolker, B.M., Brooks, M.E., Clark, C.J., Geange, S.W., Poulsen, J.R., Stevens, M.H., White, J.S., 2009. Generalized linear mixed models: a practical guide for ecology and evolution. *Trends Ecol. Evol.* 24, 127–135.
- Brand, M., Fujiwara, E., Borsutzky, S., Kalbe, E., Kessler, J., Markowitsch, H.J., 2005. Decision-making deficits of korsakoff patients in a new gambling task with explicit rules: associations with executive functions. *Neuropsychology* 19, 267.
- Breslow, N.E., Clayton, D.G., 1993. Approximate inference in Generalized Linear Mixed Models. *J. Am. Stat. Assoc.* 88, 9–25.
- Breslow, N., 2004. Whither PQL? *Proceedings of the Second Seattle Symposium in Biostatistics*. Springer, pp. 1–22.
- Cameron, A.C., Trivedi, P.K., 2013. *Regression Analysis of Count Data*. Cambridge University Press.
- Clarkson, D.B., Zhan, Y., 2002. Using spherical–radial quadrature to fit generalized linear mixed effects models. *J. Comput. Gr. Stat.* 11, 639–659.
- Crowder, M.J., 1978. Beta-binomial anova for proportions. *J. R. Stat. Soc. Ser. C (Appl. Stat.)* 27, 34–37.
- Gambarana, C., Scheggi, S., Tagliamonte, A., Tolu, P., De Montis, M.G., 2001. Animal models for the study of antidepressant activity. *Brain Res. Protoc.* 7, 11–20.
- Giang, W.C.W., Donmez, B., 2015. Interpreting visualizations of historical variability for estimating future events. In: *Proceedings of the 12th International Conference on Naturalistic Decision Making*, McLean, VA.
- Hilbe, J.M., 2013. Beta binomial regression. In: *The Selected Works of Joseph Hilbe*.
- Jaeger, T.F., 2008. Categorical data analysis: away from ANOVAs (transformation or not) and towards logit mixed models. *J. Mem. Lang.* 59, 434–446.
- James, R.J., O'Malley, C., Tunney, R.J., 2016. Why are some games more addictive than others: the effects of timing and payoff on perseverance in a slot machine game. *Front. Psychol.* 7.
- Johnson, P.C., Barry, S.J., Ferguson, H.M., Müller, P., 2015. Power analysis for generalized linear mixed models in ecology and evolution. *Methods Ecol. Evol.* 6, 133–142.
- Kain, M.P., Bolker, B.M., McCoy, M.W., 2015. A practical guide and power analysis for GLMMs: detecting among treatment variation in random effects. *PeerJ* 3, e1226.
- Kim, Y., Choi, Y.K., Emery, S., 2013. Logistic regression with multiple random effects: a simulation study of estimation methods and statistical packages. *Am. Stat.* 67.
- Kleiber, C., Zeileis, A., 2008. *Applied Econometrics with R*. Springer Science & Business Media.
- Krishnan, V., Nestler, E.J., 2008. The molecular neurobiology of depression. *Nature* 455, 894–902.
- Lazic, S.E., 2015. Analytical strategies for the marble burying test: avoiding impossible predictions and invalid p-values. *BMC Res. Notes* 8, 141.
- Lesnoff, M., Lancelot, R., 2012. aod: Analysis of overdispersed data. R package version 1.3.
- Lesnoff, M., Lancelot, R., 2013. aods3: Analysis of overdispersed data using S3 methods. aods3 Package version 0.4–1.
- Levine, R.A., Casella, G., 2001. Implementations of the Monte Carlo EM algorithm. *J. Comput. Gr. Stat.* 10, 422–439.
- Marchese, G., Scheggi, S., Secci, M.E., De Montis, M.G., Gambarana, C., 2013. Anti-anhedonic activity of long-term lithium treatment in rats exposed to repeated unavoidable stress. *Int. J. Neuropsychopharmacol. Off. Sci. J. Coll. Int. Neuropsychopharmacol.* 16, 1611–1621.
- Marszalek, J.M., Barber, C., Kohlhart, J., Holmes, C.B., 2011. Sample size in psychological research over the past 30 years. *Percept. Mot. Skills* 112, 331–348.
- Moinuddin, R., Matheson, F.I., Glazier, R.H., 2007. A simulation study of sample size for multilevel logistic regression models. *BMC Med. Res. Methodol.* 7, 34.
- Muniz-Terrera, G., van den Hout, A., Rigby, R.A., Stasinopoulos, D.M., 2016. Analysing cognitive test data: Distributions and non-parametric random effects. *Stat. Methods Med. Res.* 25 (2), 741–753.
- Paserman, M., 2016. Adaptive investment strategies during financial crises and the role of corruption: an experiment with financial professionals. Available at SSRN 2731507.
- RCoreDevelopmentTeam, 2014. *R: a Language and Environment for Statistical Computing*. R Development Core Team, Vienna, Austria.
- Raudenbush, S.W., Yang, M.-L., Yosef, M., 2000. Maximum likelihood for Generalized Linear Models with nested random effects via high-order, multivariate laplace approximation. *J. Comput. Gr. Stat.* 9, 141–157.
- Robinson, M.D., Kahraman, A., Law, C.W., Lindsay, H., Nowicka, M., Weber, L.M., Zhou, X., 2014. Statistical methods for detecting differentially methylated loci and regions. *Front. Genet.* 5.
- Scheggi, S., Pelliccia, T., Ferrari, A., De Montis, M., Gambarana, C., 2015. Imipramine, fluoxetine and clozapine differently affected reactivity to positive and negative stimuli in a model of motivational anhedonia in rats. *Neuroscience* 291, 189–202.
- Silva, J.M.C.S., Tenreiro, S., 2006. The log of gravity. *Rev. Econ. Stat.* 88, 641–658.
- Stockard, J., O'Brien, R.M., Peters, E., 2007. The use of mixed models in a modified Iowa Gambling Task and a prisoner's dilemma game. *Judgm. Decis. Mak.* 2, 9.
- van den Hout, A., Muniz-Terrera, G., 2016. Joint models for discrete longitudinal outcomes in aging research. *J. R. Stat. Soc. Ser. C (Appl. Stat.)* 65, 167–186.
- Venables, W., Ripley, B., 2002. *Modern Applied Statistics with S*.
- White, H., 1980. A heteroskedasticity-consistent covariance matrix estimator and a direct test for heteroskedasticity. *Econom. J. Econom. Soc.*, 817–838.